

CS 15 Project 4: gerp

Contents

Introduction	3
Background: grep	3
Program Design Overview	4
Files to Implement	5
Pair Programming	6
Starter Files	6
Program Specification	7
Introduction	7
Command Loop	7
Beware!	8
Output Formatting	9
What is a Word?	10
Building the Index	11
FSTree	11
DirNode	12
STL Usage	13
Compiler Options	14
Time and Space Constraints	14
Reference Implementation and Testing	15
Introduction	15
Redirecting Input/Output in unix	16
Using Diff	16
README	18
Submitting Your Work	19
Part 1: Design, FSTree, and String Manipulation	19
Design	19

FSTree Traversal	20
String Manipulation	21
Part I README	21
Provide	21
Part 2 - Final Submission	21
Helpful Tips	23
Testing	23
Hashing and Hash Tables	23
Runtime	23
Space	23
Hints	24

Introduction

We're all familiar with web search engines, and we also have tools for searching our personal computers. Have you ever wondered how the Mac Spotlight works, for example? We'll look at one approach now! In this assignment you will design and implement a program that indexes and searches files for strings. Your program will behave similarly to the **unix grep** program, which can search through all the files in a directory and look for a sequence of characters.

Background: grep

Here's an example of calling **unix grep**:

```
grep -IRn Query DirectoryToSearch
```

In this example:

- **grep** is the program we're calling
- **Query** is the target string we're searching for
- **DirectoryToSearch** is the directory where we will look for the **Query**
- the option **I** means to search all files, including binary files
- the option **R** means to traverse all files under the provided directory recursively (including following symbolic links)
- the option **n** means to print each output line from **grep** with the line number of the match of the “**Query**” within the matched file.

For example,

```
grep -IRn #include /comp/15/files
```

will produce something like

```

/comp/15/files/hw1/CharArrayList.cpp:12:#include "CharArrayList.h"
/comp/15/files/hw1/simple_exception.cpp:27:#include <iostream>
/comp/15/files/hw1/simple_exception.cpp:28:#include <string>
/comp/15/files/hw1/simple_exception.cpp:29:#include <stdexcept>
/comp/15/files/hw1/unit_tests.h:15:#include "CharArrayList.h"
/comp/15/files/hw1/unit_tests.h:16:#include <cassert>
/comp/15/files/lab2/LinkedList.cpp:13:#include "LinkedList.h"
/comp/15/files/lab2/LinkedList.cpp:14:#include <sstream>
/comp/15/files/lab2/LinkedList.cpp:15:#include <string>
/comp/15/files/lab2/LinkedList.h:15:#include "Planet.h"
/comp/15/files/lab2/Planet.cpp:11:#include "Planet.h"
/comp/15/files/lab2/Planet.cpp:12:#include <sstream>
/comp/15/files/lab2/Planet.cpp:13:#include <string>
/comp/15/files/lab2/Planet.h:14:#include <string>
/comp/15/files/lab2/planet-driver.cpp:14:#include "LinkedList.h"
/comp/15/files/lab2/planet-driver.cpp:15:#include "Planet.h"
/comp/15/files/lab2/planet-driver.cpp:16:#include <iostream>
/comp/15/files/lab2/unit_tests.h:6:#include "LinkedList.h"
/comp/15/files/lab2/unit_tests.h:7:#include "Planet.h"
/comp/15/files/lab2/unit_tests.h:8:#include <cassert>
/comp/15/files/lab2/unit_tests.h:9:#include <iostream>
/comp/15/files/hw2/CharLinkedList.cpp:12:#include "CharLinkedList.h"
/comp/15/files/lab0/welcome.cpp:7:#include <iostream>
/comp/15/files/lab1/ArrayList.cpp:15:#include "ArrayList.h"
/comp/15/files/lab1/ArrayList.cpp:16:#include <sstream>
/comp/15/files/lab1/ArrayList.h:14:#include <string>
/comp/15/files/lab1/unit_tests.h:19: *      Be sure to #include any .h files as
      necessary
/comp/15/files/lab1/unit_tests.h:37:#include "ArrayList.h"
/comp/15/files/lab1/unit_tests.h:38:#include <cassert>
/comp/15/files/lab1/unit_tests.h:39:#include <iostream>
/comp/15/files/lab1/unit_tests.h:40:#include <string>
/comp/15/files/compilation_tests/hw1_test.cpp:13:#include <iostream>
/comp/15/files/compilation_tests/hw1_test.cpp:14:#include "CharArrayList.h"
/comp/15/files/compilation_tests/hw2_test.cpp:14:#include "CharLinkedList.h"
/comp/15/files/compilation_tests/hw2_test.cpp:13:#include <iostream>

```

The first line of this output tells you that `#include` can be found on line 12 of the file located at `/comp/15/files/hw1/CharArrayList.cpp`, and that the line itself is `#include "CharArrayList.h"`. Although there are many optional flags you might use with `grep` (see <https://linux.die.net/man/1/grep> for the details), our `gerp` program is designed to replicate the core functionality above, with a twist!

Program Design Overview

How might we build something like this? We will start by limiting the constraints of our solution—whereas `grep` produces near-instantaneous results of matches at each run of the program, we will take a two-step approach with `gerp`:

- 1 Build a data structure that indexes a set of files
- 2 Use that index to respond to queries

That is, when `gerp` is run, it will require the user to provide a directory up front. Your first task will thus be to process the files of that directory (including subdirectories). You will then enter a loop where the user provides queries that your program will (quickly) answer. The queries are effectively searches for words in the indexed files. Described in English, a query might be: “Tell me all the files that have the word ‘potato’ inside the directory `/comp/15`.”

While processing the files, you might find it useful to store information about those files such as their names, their relative paths on the filesystem, and information about their contents, in one or more data structure(s) that is/are easily searched and queried. The choice of data structure(s) for this assignment is up to you!

That said, we have provided you with modules to help with indexing files (see `Building the Index` section below). Also, for this assignment, you may use certain classes from the C++ Standard Template Library (STL). Descriptions of the STL implemetations that you are allowed to use are listed in their own section named `STL Usage`.

To help you learn the interface and get a feel for the program, we have also provided you with a working reference implementation. For help on how to test your work against the reference, see the section `Testing and Reference Implementation`.

Files to Implement

We will not specify most of the files or functions you will need to write. Instead, your program is required to function as described in this specification. You may accomplish this task using any combination of files, functions, and classes you wish. We will, of course, evaluate your design.

However, in addition to writing `.h` and `.cpp` files for your classes, you will need to write a `main()` for your program, and write a `Makefile`. The default `make` action should be to compile and link the entire program, and to produce an executable program named `gerp`, which you can run by typing `./gerp` in the terminal. If you have a `clean` target in your `Makefile`, be sure you do not delete the `.o` files we give to you.

You will want to write code to test the various parts of your program separately so that you do not have to debug compound errors. To that end, you are encouraged to use the `unit_test` framework to unit test your classes; however, testing details will be up to you. As usual, you will also be required to submit a `README`. See the `README` section for details.

Pair Programming

For this project, you will be working together with a partner! You will be responsible for choosing your partner. It is essential that you and your partner have compatible working schedules, because you are required to work together for all stages of the project. For more information on our pair programming guidelines.

Starter Files

To copy the starter files, run the following command on the server

```
/comp/15/files/proj4/setup
```

Note that you should **not** run `cp`.

Program Specification

Introduction

Your program will be run from the command line like this:

```
./gerp DirectoryToIndex OutputFile
```

where `DirectoryToIndex` determines which directory will be traversed and indexed, e.g. `/comp/15/files`, and `OutputFile` names the file to which the query results will be sent.

If the user did not specify exactly two command line arguments (in addition to the program name), print this message to `std::cerr`:

```
Usage: ./gerp inputDirectory outputFile
```

and terminate the program by returning `EXIT_FAILURE` from `main`.

After being called, your program will first traverse a file tree created using a module provided by the course staff (details are in the section titled `Building the Index` below). It will index each file that it finds in the tree. After indexing all of the files, it will enter a command loop (similar to the “interactive” modes that you have implemented in previous homeworks) where the user can enter various commands to modify the search, and to quit the program.

Command Loop

Specifically, your program will print `Query?` followed by a single space to `std::cout` (**NOT** to the `OutputFile`), and then wait for a command from the user. The possible query commands are:

- `AnyString`

A word (see the `What is a Word?` section below) is treated as a query. The program will take this string and print all of the lines in the

indexed files where `AnyString` appears. Note that this is a case sensitive search, so `we` and `We` are treated as different strings/words, and so should have different results.

- `@i AnyString` or `@insensitive AnyString`
Preceding a query string by `@i` or `@insensitive` causes the program to perform a case insensitive search on the string that was passed. For example, `we` and `We` would be treated as the same string/word and will have the same results.
- `@q` or `@quit`
These commands will completely quit the program, and print:
Goodbye! Thank you and have a nice day.
This statement should be followed by a new line. Note that the program should also quit if it reaches End-Of-File (EOF).
- `@f newOutputFilename`
This command causes the program to close the current output file. Any future output should be written to the file named `newOutputFilename`.

Beware!

Beware! There are **two** output streams.

- The “Query? ” prompt always goes to `std::cout`
- The result of the query always goes to the output file, which will **either** be the `OutputFile` or the file named in the last `@f` command.

Also, treat a multi-word query as several independent 1-word queries - e.g.

```
Query? We are the champions
```

is the same as

```
Query? We
Query? are
Query? the
Query? champions
```

Note, as seen when running the reference implementation, that this will actually appear as:

```
Query? we are the champions
Query? Query? Query? Query?
```

Output Formatting

If the word (see the **What is a Word?** section below) in a query is found in the index, then, for each line it appears in, you will print to the designated output file a line of the form:

```
FileNameWithPath:LineNum: Line
```

Where:

1. **FileNameWithPath** is the full pathname of the file (including the path from the command line), followed by a colon
2. **LineNum** is the line number within that file that the query word appears on, followed by a colon and a space
3. **Line** is the full text of the line from the file
4. A newline

For example, if you ran

```
$ ./gerp small_test out.txt  
Query? we
```

which queries **we** on our **small_test** directory, and then sends output to the file **out.txt**, the file **out.txt** might read:

```
small_test/test.txt:5: we are the champions  
small_test/test.txt:6: we we we
```

NOTE: There is one newline after the last line. Also, each line that the query appears in only prints once. If the query is not found using the default search, then print:

```
query Not Found. Try with @insensitive or @i.
```

If the query is not found using the insensitive search, then print:

```
query Not Found.
```

What is a Word?

It is important to outline what a word is when dealing with a word search engine. **We will define a word as a string that starts and ends with an alphanumeric (letter or number) character.** This means that you will need to do a little string parsing to determine the output of your `gerp` implementation. To help you with this nuance we have included a couple of examples.

When searching for the word `comp` using case insensitive search, `gerp` should treat the following strings as `comp`:

- `comp`
- `comp.`
- `Comp`
- `-comp`
- `&&comp`
- `comp?!`
- `@#comp?@!`

If any of the bulleted strings were submitted as a query, `gerp` should print the lines in files that contain any of the strings on the list (however, it should print them as they exist in the file, not a processed version).

Note that `gerp` should only compare strings where all leading and trailing non-alphanumeric characters are stripped. This includes both the queries and the strings in the data files.

Note that words can contain non alphanumeric characters in the middle. For instance, `comp&!$15` is considered all one word. It should not be split into two queries as with spaces.

Building the Index

You do **NOT** need to write the **FSTree** or **DirNode** classes described below - we have implemented them for you. However, they will be critical to the success of your project!

FSTree

We will use a file-system tree to represent directories, subdirectories, and files. The data structure we will use is an **n-ary** tree, so called because a node of the tree could have any number of children. The main usage of this class is to help you navigate through folders and directories inside the computer. For example, a snapshot of a home directory might be represented in an n-ary tree like this:

```
/h/mkorman
  /coursework
    /comp11
    /comp15
      /exams
      /labs
      /assignments
      /hws
    /comp160
      /hw
  /public_html
```

Specifically, a **FSTree** is an n-ary tree which consists of **DirNodes** (which are described below). The **FSTree** class has the following public functions:

- **FSTree**(std::string rootName)
This is the constructor of the **FSTree**. It creates a file tree of **DirNodes** where the root of the tree is the directory that is passed as the parameter **rootName**. If there is an error opening directories or files, the constructor will fail and halt your program. Be careful and do not run this on just any directory - if a directory structure has a loop in it, the constructor can run forever!
- **~FSTree**()
The destructor deallocates all of the space allocated when the tree was built.
- **void burnTree**()
This function destroys the tree, frees any data that was allocated, and makes the tree empty.

- `DirNode *getRoot()`
This function returns the root of the tree. Normally, we do not want to return the private members of an object or class, however in this case it is necessary so that you can traverse the tree and index its contents.

DirNode

The `DirNode` class is the key building block of the `FSTree` class. It is our representation of a folder. Each `DirNode` instance has: a name, a list of files in the directory, and a list of subdirectories. It contains the following public methods:

- `bool hasSubDir()`
Returns true if there are any sub-directories in this directory.
- `bool hasFiles()`
Returns true if there are files in this directory.
- `bool isEmpty()`
Returns true if there are no files or sub-directories in this directory.
- `int numSubDir()`
Returns the number of sub directories in this directory.
- `int numFiles()`
Returns the number of files in this directory.
- `std::string getName()`
Returns the name of this directory.
- `DirNode *getSubDir(int n)`
Returns a pointer to the `n`th subdirectory.
- `std::string getFile(int n)`
Returns the `n`th file name in this directory.
- `DirNode *getParent()`
Get the parent directory of this directory.

The `DirNode` class also contains the following public functions, which are used to modify the contents and structure of the tree (these are used by the `FSTree` implementation). Although you have access to these functions,

you should **not** need to use them, since they may cause the tree to lose data/information. We're only telling you about them, because you'll see them listed in the `.h` file, and you may be curious.

- `DirNode(std::string newName)`
Constructor that initializes a `DirNode` with the provided folder name.
- `void setName(std::string newName)`
Sets the folder name of the current node to the provided name.
- `void addSubDirectory(DirNode *newDir)`
Adds the provided sub-directory to the current node.
- `void setParent(DirNode *newParent)`
Sets the parent node of the current node to the provided node.

In order to get a file's full path, you will need to traverse the `FSTree` and concatenate the names of the directories you encounter along the way. You will then use this full path to open the file in an `std::ifstream` and **index its contents**.

STL Usage

For this assignment, you will be allowed to use **ONLY** the following STL implementations:

- `vector`
- `queue`
- `stack`
- `set`
- `list`
- `functional`

You are not required to use any particular item of the STL. If you feel that one or more of these would be useful, you will need to learn about their respective interfaces. You can find more information about them at:

<http://www.cplusplus.com/reference>

Any other data structures that you need you must implement yourself.

Compiler Options

When compiling your implementation of `gerp`, you should compile with the flag `-O2` (That's a capital letter 'O', not the numeral zero). This will optimize your program for the system that it is compiling on, which will result in an implementation with a faster runtime. This will help during the testing phase because you will receive your results faster.

Time and Space Constraints

When designing and implementing your program you should have it build its index and run queries as quickly as possible. You may find that there is a trade-off between the two (e.g. a program that builds an index quickly may not search as fast). It is important to document your design choices, your justification of those choices, and their effects in your README.

Your program will need index our test sets and be ready to query in **under 10 minutes**. Our implementation uses approximately 2.3 GB of RAM once it has fully indexed the largest test set. The Gradescope autograder has 6GB of RAM allocated per test run. This means that you will lose points for tests if you exceed the **6GB memory requirement**. You can still get credit for tests on the smaller collections, of course, if your program works for those.

You should check the time and memory usage of your program using the following command:

```
gerp_perf [DirectoryToIndex] [OutputFile]
```

This command will display your program's memory usage and execution time of building the index. It will also display these metrics for the reference implementation. [OutputFile] will store the non-cout output of your program as usual. Note, to use this program, you will need to have added `use -q comp15` to your `.cshrc` file as you would have done to use `unit_test`.

Also, please note that our reference implementation, `the_gerp`, is by no means the fastest indexing or fastest querying solution to the problem. We hope you are able to build a faster implementation!

Reference Implementation and Testing

Introduction

In order to help you with your testing and to get familiar with the user interface expectations of this project, we have provided you with a fully compiled reference implementation called `the_gerp`. By the end of the project your `gerp` implementation should behave exactly the same as `the_gerp`. **Your version of `gerp` should behave exactly as the reference in all circumstances. Therefore, you should extensively test the reference in order to figure out how it behaves.**

Testing Directories

In order to help you test your implementations, we have provided you with some directories! The folder `/comp/15/files/proj4-test-dirs` contains all of the relevant directories, which are

- A Directory `tinyData`, which you can use to quickly test your work.
- Small, medium and large subsets of Project Gutenberg. These are samples from a massive free online library—for more information on this, see the [Project Gutenberg Website](#). Indexing/querying the medium/large subsets will give you a chance to test the time and space efficiency of your code. **This semester (Fall 2022), our autograder will not test you on the large Gutenberg dataset.** However, we will test your submission on small and medium Gutenberg, so be sure that your submission matches the expectations described earlier (indexing in under 10 minutes, no more than 6GB memory usage).

Note that there is a file named `/comp/15/files/proj4-test-dirs/dataSets.zip` - this file contains all of the testing directories, so if you'd like to download it to your system to work locally, you can do so.

Redirecting Input/Output in unix

In order to test your implementation against the reference, using redirects to send input files to and from your program and the reference will be useful. To send data **to** the `std::cin` of a unix program, do the following:

```
./programname < fileforcin.txt
```

In our case, `gerp` takes multiple arguments: a directory to query, and a file to send the program's output to. So, to send a file to the reference implementation's `std::cin`, you would do:

```
./the_gerp Directory ref_output.txt < commands.txt
```

This command will send `commands.txt` to the standard input stream for the reference implementation. You can likewise do the same with your implementation. In order to redirect a program's `std::cout` to a file, do the following:

```
./programname > fileforcin.txt
```

Again, in our case,

```
./the_gerp Directory ref_output.txt > ref_std_out.txt
```

To combine redirection of input and output, you can do the following:

```
./the_gerp Directory ref_out.txt < cmds.txt > ref_std_out.txt
```

Recall that, even though `gerp` sends **most** data to the the output file, it will send some to `std::cout` - namely, the `Query?` lines, which will then be redirected to, in this case, `ref_std_out.txt`.

Using Diff

Once you have run the same input on the reference and on your implementation, you can test that your and the reference's outputs (the contents of `ref_out.txt` above) are equivalent by:

1. Sorting the output

2. Using `diff` to compare the sorted files

Sorting is necessary because, while your program must produce output for all occurrences of the query, the order in which multiple lines appear is not specified. Print out multiple lines in whatever order your data structure and algorithm choices find convenient. To sort the output, you can use the `unix sort` command.

```
sort ref_output.txt > ref_output_sorted.txt
sort my_output.txt > my_output_sorted.txt
```

After sorting, use the `diff` command to find any differences. For example:

```
diff ref_output_sorted.txt my_output_sorted.txt
```

`diff` will print the differences in the file, if there are any. If nothing prints out, then the files are identical.

Keep in mind that, to compare the redirected `std::cout` streams of your implementation and the reference, you should **not** sort the output files before running `diff`.

README

You will be required to submit a **README** along with your code. Your **README** should have the following sections:

- A The title of the homework and the author names (you and your partner)
- B The purpose of the program
- C Acknowledgements for any help you received, including references to outside sources you consulted (though there is no need list **C++** references like `cplusplus.com`).
- D The files that you provided and a short description of what each file is and its purpose
- E How to compile and run your program
- F An “architectural overview,” i.e., a description of how your various program modules relate. For example, the **FSTree** implementation keeps a pointer to the root **DirNode**.
- G An outline of the data structures and algorithms that you used. Given that this is a data structures class, you need to always discuss any data structures that you used and justify why you used them. For this assignment it is imperative that you explain your data structures and algorithms in detail, as it will help us understand your code since there is no single right way of completing this assignment.
- H Details and an explanation of how you tested the various parts of your classes and the program as a whole. You may reference the testing files that you submitted to aid in your explanation.

Each of the sections should be clearly delineated and begin with a section heading which describes the content of the section.

Submitting Your Work

You will be submitting your work in 2 parts:

Part 1: Design, FSTree, and String Manipulation

Design

You will be required to sign up for in-person design checkoff meeting with a TA. *Both you and your partner must attend the meeting.* You can find the design checkoff form. At that meeting, among other things, you will be required to address the following items:

- What data structure(s) will you be using? Why these data structures?
CAUTION: If your answer is just vectors, you need to think again! Use the data structures we've been studying.
- What classes will you implement (presumably including one for each data structure above)? What public functions does each class support?
- How do your classes interact, i.e., does one class contain an instance of another, a pointer to an instance of another, what functions in the other class will it call?
- How will you implement insert, search?
- How do you plan to implement case-insensitive search?
- How do you plan to avoid reporting duplicate words on the same line?
- Describe the space needs of your solution and the big-O runtime of important operations (insert, search).
- Bring images of your design!

FSTree and String Manipulation

To complete this part you will need to write 2 programs:

1. a tree traversal function that prints out the full paths of each file in the tree on separate lines
2. a function that strips all leading and trailing non-alphanumeric characters from a given string.

FSTree Traversal

You should write your tree traversal function in a file named: `FSTreeTraversal.cpp`. When executed, this program should take the highest directory as a command line argument, and then print the full paths of all the files accessible from that directory:

```
./treeTraversal Directory
```

Do not worry about the order that the file paths print in, just ensure that each one of them prints. For example, if you have a directory named `Foo`, with:

- files `a.cpp`, `b.cpp`, and `c.cpp` inside it
- a subdirectory `Bar`, with:
 - files `x.cpp`, `y.cpp`, `z.cpp`

Then one possible output for running the command `./treeTraversal Foo` would be:

```
Foo/Bar/x.cpp
Foo/Bar/y.cpp
Foo/Bar/z.cpp
Foo/a.cpp
Foo/b.cpp
Foo/c.cpp
```

Again, the order here does not matter. You may assume that the directory name given to the `treeTraversal` program has no trailing `/`, e.g., the given command line argument will be `Foo` not `Foo/`.

String Manipulation

Your string processing function should be defined in a file named `stringProcessing.cpp`, and have a declaration in a file named `stringProcessing.h`. The function should have the following header:
`std::string stripNonAlphaNum (std::string input)`

The function should remove all leading and trailing non-alphanumeric characters from the given string, such that when

`@##!!#!@!#COMP-15!!!!!!` is given as a parameter the string `COMP-15` is returned.

Note: **DO NOT** provide a `main` function for this program. Feel free to write one for testing, but make sure to comment it out before submitting your work. Additionally, *you should not define this function within a class*. Just declare the function itself within the `.h`, and define it within the `.cpp`.

Part I README

Your README does not need to contain everything we normally ask for. For this phase, you can just write a very brief summary with any information you'd like the grader to see (e.g. if you have a nettlesome bug you have not yet been able to fix). Additionally, make sure to write both author's names (i.e., you and your partner).

Provide

Note: Only one of you or your partner needs to submit code. Just make sure both names are in the submitted README.

The provide command is:

```
provide comp15 proj4_gerp_phase1 FSTreeTraversal.cpp \
                                stringProcessing.h   \
                                stringProcessing.cpp \
                                README
```

Part 2 - Final Submission

For this part you will submit all of the files required (including a `Makefile`) to compile your `gerp` program. Make sure to include any testing files. Once again, only one of you or your partner should submit. The provide template is:

```
provide comp15 proj4_gerp README Makefile [Your Filenames Here]
```

Note: We cannot give you a complete provide command, so make sure you submit everything we will need to run your program. Maybe copy the files into another directoy, type **make** test the program, do a **make clean** and then provide every thing. We should be able to use **make** or **make gerp** to build your program.

Helpful Tips

Testing

We will test your solution on several directories, including `smallGutenberg` and `mediumGutenberg`. You can see some test queries and reference implementation output in `/comp/15/files/proj4-sample-execution`. You should match this output exactly (except that the order can be different as explained above).

Hashing and Hash Tables

There are example uses of the `std::hash` facility in the files we've given you. See `hash_test.cpp` and `hashExample.cpp`. If you use a hash table, then you **MUST** have it dynamically resize. You must monitor the load factor and expand if the load factor is exceeded.

Runtime

The runtime of your query processing must depend on the size of the output, not on the size of the input (that is, the index structure should be essentially **constant time** to search).

Space

Finally, space usage will be an issue. If you are not cognizant of space usage, then your program will work fine for the small or maybe even the medium collection, but will fail on the large collection when it runs out of memory. You might find the provided file `exceedmem.cpp` helpful regarding memory issues; it shows different examples of situations you might encounter when attempting to allocate too much memory.

If you are stuck, then you should get something that at least works correctly on the small data collections. To work on larger collections, you will need to choose a strategy that considers space.

Hints

- Hint: Every copy of a string requires space proportional to the length of the string. (An `std::string` is an `ArrayList` of characters.)
- Hint: **Do not store a pointer to an element in a vector.** A good final exam question would be: “Why should you not store the address of an element in a vector?” An `std::vector` is an `ArrayList`, remember.
- Hint: Pay attention to how you pass variables to functions. For example, consider the following two function signatures:
`void foo_a(string word)`
`void foo_b(string &word)`
Passing a `string` to `foo_b`, which uses a reference parameter, will be more space and time efficient than passing a `string` to `foo_a`. See why? Leverage this to your advantage!